



```
from google.colab import files
uploaded = files.upload()
```

 No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Savino CarPrice Assignment.csv to CarPrice Assignment.csv

```
import pandas as pd
```


```
df = pd.read_csv('CarPrice_Assignment.csv')
df.head()
```



	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	engine location	wheelbase	...	en
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	front	88.6	...	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	front	88.6	...	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	...	
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	front	99.8	...	
4	5	2	audi 100ls	gas	std	four	sedan	4wd	front	99.4	...	

5 rows x 26 columns

```
df = pd.read_csv('CarPrice_Assignment.csv')
df.head()
```



	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	engine location	wheelbase	...	en
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	front	88.6	...	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	front	88.6	...	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	...	
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	front	99.8	...	
4	5	2	audi 100ls	gas	std	four	sedan	4wd	front	99.4	...	

5 rows x 26 columns


```
#Data Cleaning
df.drop(['car_ID'], axis=1, inplace=True)

df['CarCompany'] = df['CarName'].apply(lambda x: x.split(' ')[0].lower())

df.drop(['CarName'], axis=1, inplace=True)

df['CarCompany'] = df['CarCompany'].replace({
    'vw': 'volkswagen',
    'vokswagen': 'volkswagen',
    'porcshce': 'porsche',
    'toyouta': 'toyota',
    'maxda': 'mazda',
    'Nissan': 'nissan'
})

df[['CarCompany']].drop_duplicates().sort_values(by='CarCompany')
```

 **CarCompany**

0	alfa-romero
3	audi
10	bmw
67	buick
18	chevrolet
21	dodge
30	honda
43	isuzu
47	jaguar
50	mazda
75	mercury
76	mitsubishi
89	nissan
107	peugeot
118	plymouth
125	porsche
130	renault
132	saab
138	subaru
150	toyota
182	volkswagen
194	volvo


```
import matplotlib.pyplot as plt
import seaborn as sns

numeric_df = df.select_dtypes(include=['number'])

plt.figure(figsize=(12, 8))
sns.heatmap(numeric_df.corr(), annot=True, cmap='coolwarm')
plt.title("Heatmap of Numeric Feature Correlations")
plt.show()
```



Heatmap of Numeric Feature Correlations



```
df.drop(['carlength', 'carwidth', 'curbweight', 'highwaympg'], axis=1, inplace=True)
```

```
df.columns
```

```
Index(['symboling', 'fueltype', 'aspiration', 'doornumber', 'carbody',
       'drivewheel', 'enginelocation', 'wheelbase', 'carheight', 'enginetype',
       'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke',
       'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'price',
       'CarCompany'],
      dtype='object')
```

```
df = pd.get_dummies(df, drop_first=True)
```

```
df.head()
```

```

symboling  wheelbase  carheight  enginesize  boreratio  stroke  compressionratio  horsepower  peakrpm  citympg  ...  Car
0          3        88.6       48.8        130        3.47    2.68                9.0         111    5000    21    ...
1          3        88.6       48.8        130        3.47    2.68                9.0         111    5000    21    ...
2          1        94.5       52.4        152        2.68    3.47                9.0         154    5000    19    ...
3          2        99.8       54.3        109        3.19    3.40               10.0         102    5500    24    ...
4          2        99.4       54.3        136        3.19    3.40                8.0         115    5500    18    ...

```

```
5 rows x 61 columns
```

```
from sklearn.model_selection import train_test_split
```

```
X = df.drop('price', axis=1)
y = df['price']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.7, random_state=100)
```

```
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
((143, 60), (62, 60), (143,), (62,))
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler()
```

```
X_train_scaled = pd.DataFrame(scaler.fit_transform(X_train), columns=X_train.columns)
```

```
X_test_scaled = pd.DataFrame(scaler.transform(X_test), columns=X_test.columns)
```

```
X_train_scaled.head()
```

```

symboling  wheelbase  carheight  enginesize  boreratio  stroke  compressionratio  horsepower  peakrpm  citympg  ...  Ca
0          0.6    0.244828   0.265487   0.139623   0.230159  0.525253           0.15000    0.083333  0.551020  0.500000  ...
1          1.0    0.272414   0.212389   0.339623   1.000000  0.464646           0.15625    0.395833  0.551020  0.166667  ...
2          0.6    0.272414   0.424779   0.139623   0.444444  0.449495           0.15000    0.266667  1.000000  0.361111  ...
3          1.0    0.068966   0.088496   0.260377   0.626984  0.247475           0.12500    0.262500  0.346939  0.222222  ...
4          0.2    0.610345   0.858407   0.260377   0.746032  0.484848           0.03125    0.475000  0.387755  0.111111  ...
5 rows x 60 columns

```

```

from sklearn.linear_model import LinearRegression
from sklearn.feature_selection import RFE

```

```
lm = LinearRegression()
```

```
rfe = RFE(estimator=lm, n_features_to_select=15)
```

```
rfe = rfe.fit(X_train_scaled, y_train)
```

```

selected_cols = X_train_scaled.columns[rfe.support_]
selected_cols

```

```

Index(['wheelbase', 'enginesize', 'boreratio', 'stroke', 'engineloation_rear',
      'enginetype_rotor', 'cylindernumber_five', 'cylindernumber_four',
      'cylindernumber_three', 'cylindernumber_twelve', 'cylindernumber_two',
      'CarCompany_bmw', 'CarCompany_peugeot', 'CarCompany_porsche',
      'CarCompany_saab'],
      dtype='object')

```

```
X_train_rfe = X_train_scaled[selected_cols]
```

```
y_train = y_train.reset_index(drop=True)
```

```
import statsmodels.api as sm
```

```
X_train_sm = sm.add_constant(X_train_rfe)
```

```
model = sm.OLS(y_train, X_train_sm).fit()
```

```
print(model.summary())
```

```

OLS Regression Results
=====
Dep. Variable:          price      R-squared:          0.920
Model:                  OLS        Adj. R-squared:       0.911
Method:                 Least Squares      F-statistic:       105.2
Date:                   Sat, 21 Jun 2025    Prob (F-statistic): 4.83e-63
Time:                   09:25:23          Log-Likelihood:   -1303.4
No. Observations:       143              AIC:             2637.
Df Residuals:           128              BIC:             2681.
Df Model:                14
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-2681.9757	1522.550	-1.762	0.081	-5694.601	330.649
wheelbase	9345.5364	1728.433	5.407	0.000	5925.536	1.28e+04
enginesize	6.605e+04	5717.693	11.552	0.000	5.47e+04	7.74e+04
boreratio	-1.224e+04	2243.488	-5.457	0.000	-1.67e+04	-7804.384
stroke	-1.063e+04	1935.359	-5.494	0.000	-1.45e+04	-6803.369
engineloation_rear	6332.9278	2917.617	2.171	0.032	559.924	1.21e+04
enginetype_rotor	1.103e+04	1354.568	8.145	0.000	8353.331	1.37e+04
cylindernumber_five	8604.9276	1291.905	6.661	0.000	6048.672	1.12e+04
cylindernumber_four	8165.8463	1618.410	5.046	0.000	4963.545	1.14e+04
cylindernumber_three	1.4e+04	3068.240	4.562	0.000	7927.684	2.01e+04
cylindernumber_twelve	-2.092e+04	3955.073	-5.288	0.000	-2.87e+04	-1.31e+04
cylindernumber_two	1.103e+04	1354.568	8.145	0.000	8353.331	1.37e+04

```

CarCompany_bmw      8500.8200   1091.750    7.786    0.000   6340.606   1.07e+04
CarCompany_peugeot  -2178.9655   1119.054   -1.947    0.054  -4393.206    35.275
CarCompany_porsche    1.04e+04   1843.467    5.644    0.000   6756.984   1.41e+04
CarCompany_saab       3941.6876   1391.261    2.833    0.005   1188.840   6694.535
=====
Omnibus:                20.586   Durbin-Watson:                1.912
Prob(Omnibus):          0.000   Jarque-Bera (JB):              29.137
Skew:                   0.780   Prob(JB):                      4.71e-07
Kurtosis:               4.567   Cond. No.                      1.64e+17
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The smallest eigenvalue is 1.25e-32. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
X_train_rfe = X_train_rfe.drop('CarCompany_peugeot', axis=1)
```

```

X_train_sm = sm.add_constant(X_train_rfe)
model = sm.OLS(y_train, X_train_sm).fit()
print(model.summary())

```



OLS Regression Results

```

=====
Dep. Variable:          price   R-squared:                0.918
Model:                  OLS    Adj. R-squared:            0.909
Method:                 Least Squares   F-statistic:              110.6
Date:                   Sat, 21 Jun 2025   Prob (F-statistic):       2.87e-63
Time:                   09:27:14   Log-Likelihood:           -1305.5
No. Observations:       143   AIC:                      2639.
Df Residuals:           129   BIC:                      2681.
Df Model:                13
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-1807.3445	1470.437	-1.229	0.221	-4716.641	1101.952
wheelbase	7423.9684	1434.279	5.176	0.000	4586.212	1.03e+04
engineize	6.581e+04	5777.814	11.389	0.000	5.44e+04	7.72e+04
boreratio	-1.202e+04	2264.580	-5.306	0.000	-1.65e+04	-7536.458
stroke	-1.025e+04	1946.119	-5.267	0.000	-1.41e+04	-6400.546
enginelocation_rear	5662.5883	2928.413	1.934	0.055	-131.349	1.15e+04
enginetype_rotor	1.073e+04	1359.969	7.889	0.000	8037.985	1.34e+04
cylindernumber_five	8680.7402	1305.215	6.651	0.000	6098.340	1.13e+04
cylindernumber_four	7561.7451	1605.488	4.710	0.000	4385.248	1.07e+04
cylindernumber_three	1.304e+04	3061.069	4.260	0.000	6983.628	1.91e+04
cylindernumber_twelve	-2.079e+04	3997.090	-5.201	0.000	-2.87e+04	-1.29e+04
cylindernumber_two	1.073e+04	1359.969	7.889	0.000	8037.985	1.34e+04
CarCompany_bmw	8612.3999	1101.978	7.815	0.000	6432.110	1.08e+04
CarCompany_porsche	1.019e+04	1859.922	5.478	0.000	6508.483	1.39e+04
CarCompany_saab	4230.7827	1398.202	3.026	0.003	1464.405	6997.161

```

=====
Omnibus:                22.726   Durbin-Watson:                1.913
Prob(Omnibus):          0.000   Jarque-Bera (JB):              35.142
Skew:                   0.808   Prob(JB):                      2.34e-08
Kurtosis:               4.813   Cond. No.                      1.39e+17
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The smallest eigenvalue is 1.75e-32. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
y_train_pred = model.predict(X_train_sm)
```

```
residuals = y_train - y_train_pred
```

```

import seaborn as sns
import matplotlib.pyplot as plt

```

```

sns.histplot(residuals, kde=True)
plt.title("Residuals Distribution")
plt.xlabel("Error")
plt.ylabel("Frequency")
plt.show()

```

